

# [Summary] Unsupervised Learning

## Clustering

### Unsupervised Learning: Introduction

Unsupervised learning is contrasted from supervised learning because it uses an unlabeled training set rather than a labeled one.

Clustering is good for:

- Market segmentation
- Social network analysis
- Organizing computer clusters
- Astronomical data analysis

### K-Means Algorithm

The K-Means Algorithm is the most popular and widely used algorithm for automatically grouping data into coherent subsets.

1. Randomly initialize two points in the dataset called the cluster centroids.
2. Cluster assignment: assign all examples into one of two groups based on which cluster centroid the example is closest to.

3. Move centroid: compute the averages for all the points inside each of the two cluster centroid groups, then move the cluster centroid points to those averages.
4. Re-run (2) and (3) until we have found our clusters.

Our main variables are:

- K (number of clusters)
- Training set  $x^1, x^2, \dots, x^m$
- Where  $x^i \in \mathbb{R}^n$

Note that we **will not use** the  $x_0 = 1$  convention.

## Cluster Assignment

The first for-loop is the 'Cluster Assignment' step. We make a vector  $c$  where  $c(i)$  represents the centroid assigned to example  $x(i)$ .

We can write the operation of the Cluster Assignment step more mathematically as follows:

$$c^i = \arg \min_k \|x^i - \mu_k\|^2$$

That is, each  $c^i$  contains the index of the centroid that has minimal distance to  $x^i$ .

The square convention serves two purposes, **minimize more sharply** and **less computation**.

## Move Centroid

The **second for-loop** is the '**Move Centroid**' step where we **move** each centroid **to** the average of its group.

More formally, the equation for this loop is as follows:

$$\mu_k = \frac{1}{n} [x^{k_1} + x^{k_2} + \dots + x^{k_n}] \in \mathbb{R}^n$$

Where each of  $x^{k_1}, x^{k_2}, \dots, x^{k_n}$  are the training examples assigned to group  $m_{\mu_k}$

If you have a cluster centroid with **0 points** assigned to it, you can **randomly re-initialize** that centroid to a new point. You can also simply **eliminate** that cluster group.

After a number of iterations the algorithm will **converge**, where new iterations do not affect the clusters.

Note on non-separated clusters: some datasets have no real inner separation or natural structure. K-means can still evenly segment your data into K subsets, so can still be useful in this case.

## Optimization Objective

Recall some of the parameters we used in our algorithm:

- $c^i$  = index of cluster (1,2,...,K) to which example  $x(i)$  is currently assigned
- $\mu_k$  = cluster centroid  $k$  ( $\mu_k \in \mathbb{R}^n$ )

- $\mu_{c^i}$  = cluster centroid of cluster to which example  $x(i)$  has been assigned

Using these variables we can define our **cost function**:

$$J(c^1, \dots, c^m, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m \|x^i - \mu_{c^i}\|^2$$

Our **optimization objective** is to minimize all our parameters using the above cost function:

$$\min_{c, \mu} J(c, \mu)$$

That is, we are finding all the values in sets  $c$ , representing all our clusters, and  $\mu$ , representing all our centroids, that will minimize **the average of the distances** of every training example to its corresponding cluster centroid.

The above cost function is often called the **distortion** of the training examples.

## Cluster assignment

In the **cluster assignment step**, our goal is to:

Minimize  $J(\dots)$  with  $c^1, \dots, c^m$  (holding  $\mu_1, \dots, \mu_K$  fixed)

In the **move centroid step**, our goal is to:

Minimize  $J(\dots)$  with  $\mu_1, \dots, \mu_K$

With k-means, it is **not possible for the cost function to sometimes increase**. It should always descend.

## Random Initialization

There's one particular recommended method for randomly initializing your cluster centroids.

1. Have  $K < m$ . That is, make sure the number of your clusters is less than the number of your training examples.
2. Randomly **pick K training examples**. (the selected examples are **unique**).
3. Set  $\mu_1, \dots, \mu_K$  equal to these K examples.

K-means **can get stuck in local optima**. To decrease the chance of this happening, you can run the algorithm on **many different random initializations**. In cases where  $K < 10$  it is strongly recommended to run a loop of random initializations.

## Choosing the Number of Clusters

Choosing K can be quite arbitrary and ambiguous.

**The elbow method**: plot the cost J and the number of clusters K. The cost function should reduce as we increase the number of clusters, and then flatten out. Choose K at the point where the cost function starts to flatten out.

However, fairly often, the curve is **very gradual**, so there's no clear elbow.

**Note:** J will **always** decrease as K is increased. The one exception is if k-means gets stuck at a bad local optimum.

Another way to choose K is to observe how well k-means performs on a **downstream purpose**. In other words, you choose K that proves to be most useful for some goal you're trying to achieve from using these clusters.